

Wrangling data with dplyr

2019-08-15

dplyr : go wrangling



The main verbs of dplyr

`select()`

`filter()`

`mutate()`

`arrange()`

`summarize()`

`group_by()`



The main verbs of dplyr

select() = **Subset columns (variables)**

filter()

mutate()

arrange()

summarize()

group_by()

select()

```
select(<DATA>, <VARIABLES>)
```

select()

```
select(<DATA>, <VARIABLES>)
```

```
diamonds
```

```
## # A tibble: 53,940 x 10
```

##		carat	cut	color	clarity	depth	table	price	x	y
##		<dbl>	<ord>	<ord>	<ord>	<dbl>	<dbl>	<int>	<dbl>	<dbl>
##	1	0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98
##	2	0.21	Premium	E	SI1	59.8	61	326	3.89	3.84
##	3	0.23	Good	E	VS1	56.9	65	327	4.05	4.07
##	4	0.29	Premium	I	VS2	62.4	58	334	4.2	4.23
##	5	0.31	Good	J	SI2	63.3	58	335	4.34	4.35
##	6	0.24	Very G...	J	VVS2	62.8	57	336	3.94	3.96
##	7	0.24	Very G...	I	VVS1	62.3	57	336	3.95	3.98
##	8	0.26	Very G...	H	SI1	61.9	55	337	4.07	4.11
##	9	0.22	Fair	E	VS2	65.1	61	337	3.87	3.78
##	10	0.23	Very G...	H	VS1	59.4	61	338	4	4.05

```
## # ... with 53,930 more rows, and 1 more variable: z <dbl>
```



new data alert!



diamonds

	carat	cut	color	clarity	depth	table	price	x	y	z
1	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
2	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
3	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
4	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
5	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75
6	0.24	Very Good	J	VVS2	62.8	57.0	336	3.94	3.96	2.48
7	0.24	Very Good	I	VVS1	62.3	57.0	336	3.95	3.98	2.47
8	0.26	Very Good	H	SI1	61.9	55.0	337	4.07	4.11	2.53
9	0.22	Fair	E	VS2	65.1	61.0	337	3.87	3.78	2.49
10	0.23	Very Good	H	VS1	59.4	61.0	338	4.00	4.05	2.39
11	0.30	Good	J	SI1	64.0	55.0	339	4.25	4.28	2.73
12	0.23	Ideal	J	VS1	62.8	56.0	340	3.93	3.90	2.46
13	0.22	Premium	F	SI1	60.4	61.0	342	3.88	3.84	2.33
14	0.31	Ideal	J	SI2	62.2	54.0	344	4.35	4.37	2.71
15	0.20	Premium	E	SI2	60.2	62.0	345	3.79	3.75	2.27
16	0.32	Premium	E	I1	60.9	58.0	345	4.38	4.42	2.68
17	0.30	Ideal	I	SI2	62.0	54.0	348	4.31	4.34	2.68
18	0.30	Good	J	SI1	63.4	54.0	351	4.23	4.29	2.70

Where does it come from?

The `ggplot2` R package

How can I use it?

```
library(ggplot2)
View(diamonds)
```



it's invisible!

select()

```
select(diamonds, carat, cut, color, clarity)
```


select()

```
select(diamonds, carat, cut, color, clarity)
```

```
## # A tibble: 53,940 x 4
```

```
##   carat cut      color clarity
```

```
##   <dbl> <ord>    <ord> <ord>
```

```
## 1 0.23 Ideal      E      SI2
```

```
## 2 0.21 Premium    E      SI1
```

```
## 3 0.23 Good       E      VS1
```

```
## 4 0.290 Premium   I      VS2
```

```
## 5 0.31 Good       J      SI2
```

```
## 6 0.24 Very Good  J      VVS2
```

```
## 7 0.24 Very Good  I      VVS1
```

```
## 8 0.26 Very Good  H      SI1
```

```
## 9 0.22 Fair       E      VS2
```

```
## 10 0.23 Very Good  H      VS1
```

```
## # ... with 53,930 more rows
```

select()

```
select(diamonds, carat, cut, color, clarity)
```

```
select(diamonds, carat:clarity)
```

```
select(diamonds, 1:4)
```

```
select(diamonds, starts_with("c"))
```

```
?select_helpers
```

gapminder

```
library(gapminder)
gapminder
```

```
## # A tibble: 1,704 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```



new data alert!



gapminder

	country	continent	year	lifeExp	pop	gdpPercap
1	Afghanistan	Asia	1952	28.801	8425333	779.4453
2	Afghanistan	Asia	1957	30.332	9240934	820.8530
3	Afghanistan	Asia	1962	31.997	10267083	853.1007
4	Afghanistan	Asia	1967	34.020	11537966	836.1971
5	Afghanistan	Asia	1972	36.088	13079460	739.9811
6	Afghanistan	Asia	1977	38.438	14880372	786.1134
7	Afghanistan	Asia	1982	39.854	12881816	978.0114
8	Afghanistan	Asia	1987	40.822	13867957	852.3959
9	Afghanistan	Asia	1992	41.674	16317921	649.3414
10	Afghanistan	Asia	1997	41.763	22227415	635.3414
11	Afghanistan	Asia	2002	42.129	25268405	726.7341
12	Afghanistan	Asia	2007	43.828	31889923	974.5803
13	Albania	Europe	1952	55.230	1282697	1601.0561
14	Albania	Europe	1957	59.280	1476505	1942.2842
15	Albania	Europe	1962	64.820	1728137	2312.8890

Where does it come from?

The gapminder R
package

How can I use it?

```
library(gapminder)  
View(gapminder)
```



it's invisible!

Your turn 1

Alter the code to select just the pop column:

```
select(gapminder, year, lifeExp)
```

Your Turn 1

```
select(gapminder, pop)
```

```
## # A tibble: 1,704 x 1
##       pop
##   <int>
## 1  8425333
## 2  9240934
## 3 10267083
## 4 11537966
## 5 13079460
## 6 14880372
## 7 12881816
## 8 13867957
## 9 16317921
## 10 22227415
## # ... with 1,694 more rows
```

Show of Hands

Which of these is NOT a way to select the country and continent columns together?

```
select(gapminder, -c(year, lifeExp, pop, gdpPercap))
```

```
select(gapminder, country:continent)
```

```
select(gapminder, starts_with("c"))
```

```
select(gapminder, ends_with("t"))
```

Show of Hands

Which of these is NOT a way to select the country and continent columns together?

```
select(gapminder, ends_with("t"))
```

```
## # A tibble: 1,704 x 1
```

```
##   continent
```

```
##   <fct>
```

```
## 1 Asia
```

```
## 2 Asia
```

```
## 3 Asia
```

```
## 4 Asia
```

```
## 5 Asia
```

```
## 6 Asia
```

```
## 7 Asia
```

```
## 8 Asia
```

```
## 9 Asia
```

```
## 10 Asia
```

```
## # ... with 1,694 more rows
```


The main verbs of dplyr

`select()`

`filter()` = **Subset rows by value**

`mutate()`

`arrange()`

`summarize()`

`group_by()`

filter()

```
filter(<DATA>, <PREDICATES>)
```

Predicates: TRUE or FALSE statements

filter()

```
filter(<DATA>, <PREDICATES>)
```

Predicates: TRUE or FALSE statements

Comparisons: **>**, **>=**, **<**, **<=**, **!=** (not equal), and **==** (equal).

filter()

```
filter(<DATA>, <PREDICATES>)
```

Predicates: TRUE or FALSE statements

Comparisons: >, >=, <, <=, != (not equal), and == (equal).

Operators: & is "and", | is "or", and ! is "not"

filter()

```
filter(<DATA>, <PREDICATES>)
```

Predicates: TRUE or FALSE statements

Comparisons: >, >=, <, <=, != (not equal), and == (equal).

Operators: & is "and", | is "or", and ! is "not"

%in%

```
"a" %in% c("a", "b", "c")
```

```
## [1] TRUE
```

filter()

```
filter(diamonds, cut == "Ideal", carat > 3)
```

filter()

```
filter(diamonds, cut == "Ideal", carat > 3)
```

```
## # A tibble: 4 x 10
```

```
##   carat cut    color clarity depth table price      x      y
##   <dbl> <ord> <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl>
```

```
## 1  3.22 Ideal I      I1      62.6    55 12545  9.49  9.42
```

```
## 2  3.5  Ideal H      I1      62.8    57 12587  9.65  9.59
```

```
## 3  3.01 Ideal J      SI2     61.7    58 16037  9.25  9.2
```

```
## 4  3.01 Ideal J      I1      65.4    60 16538  8.99  8.93
```

```
## # ... with 1 more variable: z <dbl>
```

Your turn 2

Show:

All of the rows where pop is greater than or equal to 100000

All of the rows for El Salvador

**All of the rows that have a missing value for year
(no need to edit this code)**

Your turn 2

Show:

All of the rows where pop is greater than or equal to 100000

All of the rows for El Salvador

**All of the rows that have a missing value for year
(no need to edit this code)**

```
filter(gapminder, pop >= 100000)
filter(gapminder, country == "El Salvador")
filter(gapminder, is.na(year))
```

filter()

```
filter(diamonds, cut == "Ideal" | cut == "Very Good", carat > 3)
```

```
## # A tibble: 6 x 10
```

```
##   carat cut      color clarity depth table price      x      y
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl>
## 1  3.22 Ideal    I      I1      62.6   55 12545  9.49  9.42
## 2  3.5  Ideal    H      I1      62.8   57 12587  9.65  9.59
## 3  3.04 Very Go... I      SI2     63.2   59 15354  9.14  9.07
## 4  4     Very Go... I      I1      63.3   58 15984 10.0   9.94
## 5  3.01 Ideal    J      SI2     61.7   58 16037  9.25  9.2
## 6  3.01 Ideal    J      I1      65.4   60 16538  8.99  8.93
## # ... with 1 more variable: z <dbl>
```

Your turn 3

Use Boolean operators to alter the code below to return only the rows that contain:

El Salvador

Countries that had populations over 100000 in 1960 or earlier

```
filter(gapminder, country == "El Salvador" | country == "Oman")  
filter(_____, _____)
```

Your turn 3

Use Boolean operators to alter the code below to return only the rows that contain:

El Salvador

Countries that had populations over 100000 in 1960 or earlier

```
filter(gapminder, country == "El Salvador")  
filter(gapminder, pop > 100000, year <= 1960)
```

The main verbs of dplyr

`select()`

`filter()`

`mutate()` = **Change or add a variable**

`arrange()`

`summarize()`

`group_by()`

mutate()

```
mutate(<DATA>, <NAME> = <FUNCTION>)
```

mutate()

```
mutate(diamonds, log_price = log(price), log_pricesq = log_price^2)
```

mutate()

```
mutate(diamonds, log_price = log(price), log_pricesq = log_price^2)
```

```
## # A tibble: 53,940 x 12
```

```
##   carat cut      color clarity depth table price      x      y
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl>
## 1 0.23 Ideal    E      SI2      61.5    55    326    3.95    3.98
## 2 0.21 Premium E      SI1      59.8    61    326    3.89    3.84
## 3 0.23 Good     E      VS1      56.9    65    327    4.05    4.07
## 4 0.290 Premium I      VS2      62.4    58    334    4.2     4.23
## 5 0.31 Good     J      SI2      63.3    58    335    4.34    4.35
## 6 0.24 Very G... J      VVS2      62.8    57    336    3.94    3.96
## 7 0.24 Very G... I      VVS1      62.3    57    336    3.95    3.98
## 8 0.26 Very G... H      SI1      61.9    55    337    4.07    4.11
## 9 0.22 Fair     E      VS2      65.1    61    337    3.87    3.78
## 10 0.23 Very G... H      VS1      59.4    61    338    4       4.05
## # ... with 53,930 more rows, and 3 more variables: z <dbl>,
## #   log_price <dbl>, log_pricesq <dbl>
```


The main verbs of dplyr

`select()`

`filter()`

`mutate()`

`arrange()` = **Sort the data set**

`summarize()`

`group_by()`

arrange()

```
arrange(<DATA>, <SORTING VARIABLE>)
```

arrange()

```
arrange(diamonds, price)
```

```
## # A tibble: 53,940 x 10
```

```
##   carat cut      color clarity depth table price      x      y
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl>
## 1 0.23 Ideal    E      SI2      61.5    55    326    3.95    3.98
## 2 0.21 Premium E      SI1      59.8    61    326    3.89    3.84
## 3 0.23 Good     E      VS1      56.9    65    327    4.05    4.07
## 4 0.290 Premium I      VS2      62.4    58    334    4.2     4.23
## 5 0.31 Good     J      SI2      63.3    58    335    4.34    4.35
## 6 0.24 Very G... J      VVS2      62.8    57    336    3.94    3.96
## 7 0.24 Very G... I      VVS1      62.3    57    336    3.95    3.98
## 8 0.26 Very G... H      SI1      61.9    55    337    4.07    4.11
## 9 0.22 Fair     E      VS2      65.1    61    337    3.87    3.78
## 10 0.23 Very G... H      VS1      59.4    61    338    4       4.05
## # ... with 53,930 more rows, and 1 more variable: z <dbl>
```

arrange()

```
arrange(diamonds, cut, price)
```

```
## # A tibble: 53,940 x 10
##   carat cut    color clarity depth table price      x      y
##   <dbl> <ord> <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl>
## 1  0.22 Fair    E      VS2     65.1    61   337   3.87   3.78
## 2  0.25 Fair    E      VS1     55.2    64   361   4.21   4.23
## 3  0.23 Fair    G      VVS2    61.4    66   369   3.87   3.91
## 4  0.27 Fair    E      VS1     66.4    58   371   3.99   4.02
## 5  0.3   Fair    J      VS2     64.8    58   416   4.24   4.16
## 6  0.3   Fair    F      SI1     63.1    58   496   4.3    4.22
## 7  0.34 Fair    J      SI1     64.5    57   497   4.38   4.36
## 8  0.37 Fair    F      SI1     65.3    56   527   4.53   4.47
## 9  0.3   Fair    D      SI2     64.6    54   536   4.29   4.25
## 10 0.25 Fair    D      VS1     61.2    55   563   4.09   4.11
## # ... with 53,930 more rows, and 1 more variable: z <dbl>
```

desc()

```
arrange(diamonds, cut, desc(price))
```

```
## # A tibble: 53,940 x 10
##   carat cut    color clarity depth table price      x      y
##   <dbl> <ord> <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl>
## 1  2.01 Fair    G      SI1     70.6    64 18574  7.43  6.64
## 2  2.02 Fair    H      VS2     64.5    57 18565   8    7.95
## 3  4.5  Fair    J      I1     65.8    58 18531 10.2  10.2
## 4  2    Fair    G      VS2     67.6    58 18515  7.65  7.61
## 5  2.51 Fair    H      SI2     64.7    57 18308  8.44  8.5
## 6  3.01 Fair    I      SI2     65.8    56 18242  8.99  8.94
## 7  3.01 Fair    I      SI2     65.8    56 18242  8.99  8.94
## 8  2.32 Fair    H      SI1     62     62 18026  8.47  8.31
## 9  5.01 Fair    J      I1     65.5    59 18018 10.7  10.5
## 10 1.93 Fair    F      VS1     58.9    62 17995  8.17  7.97
## # ... with 53,930 more rows, and 1 more variable: z <dbl>
```

Your turn 4

Arrange gapminder by year. Add lifeExp as a second (tie breaking) variable to arrange on.

Which country had the lowest life expectancy in 1952?

Your turn 4

```
arrange(gapminder, year, lifeExp)
```

```
## # A tibble: 1,704 x 6
```

```
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>   <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Gambia       Africa    1952   30    284320    485.
## 3 Angola       Africa    1952  30.0  4232095   3521.
## 4 Sierra Leone Africa    1952  30.3  2143249    880.
## 5 Mozambique   Africa    1952  31.3  6446316    469.
## 6 Burkina Faso Africa    1952  32.0  4469979    543.
## 7 Guinea-Bissau Africa    1952  32.5   580653    300.
## 8 Yemen, Rep.  Asia      1952  32.5  4963829    782.
## 9 Somalia      Africa    1952  33.0  2526994   1136.
## 10 Guinea      Africa    1952  33.6  2664249    510.
## # ... with 1,694 more rows
```

Your turn 5

Use `desc()` to find the country with the highest `gdpPercap`.

Your turn 5

```
arrange(gapminder, desc(gdpPercap))
```

```
## # A tibble: 1,704 x 6
```

```
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>      <int>   <dbl>    <int>    <dbl>
## 1 Kuwait      Asia        1957    58.0   212846   113523.
## 2 Kuwait      Asia        1972    67.7   841934   109348.
## 3 Kuwait      Asia        1952    55.6   160000   108382.
## 4 Kuwait      Asia        1962    60.5   358266    95458.
## 5 Kuwait      Asia        1967    64.6   575003    80895.
## 6 Kuwait      Asia        1977    69.3  1140357    59265.
## 7 Norway      Europe       2007    80.2  4627926   49357.
## 8 Kuwait      Asia        2007    77.6  2505559   47307.
## 9 Singapore   Asia        2007    80.0  4553009   47143.
## 10 Norway     Europe       2002    79.0  4535591   44684.
## # ... with 1,694 more rows
```

Detour: The Pipe

%>%

**Passes the result on one function to
another function**

Detour: The Pipe

```
diamonds <- arrange(diamonds, price)
diamonds <- filter(diamonds, price > 300)
diamonds <- mutate(diamonds, log_price = log(price))

diamonds
```

Detour: The Pipe

```
diamonds <- diamonds %>%  
  arrange(price) %>%  
  filter(price > 300) %>%  
  mutate(log_price = log(price))  
  
diamonds
```

Keyboard shortcuts

Insert <- with alt/opt + -

Insert %>% with ctrl/cmd + shift + m

Your turn 6

Use %>% to write a sequence of functions that:

- 1. Filter only countries that are in the continent of Oceania.**
- 2. Select the country, year and lifeExp columns**
- 3. Arrange the results so that the highest life expectancy is at the top.**

Your turn 6

```
gapminder %>%  
  filter(continent == "Oceania") %>%  
  select(country, year, lifeExp) %>%  
  arrange(desc(lifeExp))
```

```
## # A tibble: 24 x 3  
##   country      year lifeExp  
##   <fct>      <int>   <dbl>  
## 1 Australia    2007    81.2  
## 2 Australia    2002    80.4  
## 3 New Zealand  2007    80.2  
## 4 New Zealand  2002    79.1  
## 5 Australia    1997    78.8  
## 6 Australia    1992    77.6  
## 7 New Zealand  1997    77.6  
## 8 New Zealand  1992    76.3  
## 9 Australia    1987    76.3  
## 10 Australia   1982    74.7  
## # ... with 14 more rows
```

Challenge!

- 1. Import the diabetes data from the importing data. A copy of the CSV file is available in this folder.**
- 2. Add the variable bmi to the data set using height and weight using the formula: $(\text{weight} / \text{height}^2) * 703$**
- 3. Select just id, glyhgb, and the new variable you created.**
- 4. Filter rows that have BMI > 35. How many rows and columns are in your new data set?**


```
diabetes <- read_csv("diabetes.csv")
diabetes %>%
  mutate(bmi = (weight / height^2) * 703) %>%
  select(id, glyhb, bmi) %>%
  filter(bmi > 35)
```

```
diabetes <- read_csv("diabetes.csv")
diabetes %>%
  mutate(bmi = (weight / height^2) * 703) %>%
  select(id, glyhb, bmi) %>%
  filter(bmi > 35)
```

```
## # A tibble: 61 x 3
```

```
##       id glyhb  bmi
##   <dbl> <dbl> <dbl>
## 1  1001  4.44  37.4
## 2  1002  4.64  48.4
## 3  1022  5.78  35.8
## 4  1029  4.97  40.8
## 5  1253  4.67  36.0
## 6  1254 12.7   42.5
## 7  1280  5.10  38.3
## 8  1501  4.41  40.0
## 9  2753  5.57  35.3
## 10 2757  6.33  35.3
## # ... with 51 more rows
```

The main verbs of dplyr

`select()`

`filter()`

`mutate()`

`arrange()`

`summarize()` = **Summarize the data**

`group_by()` = **Group the data**

summarize()

```
summarize(<DATA>, <NAME> = <FUNCTION>)
```

summarize()

```
summarize(diamonds, n = n(), mean_price = mean(price))
```

```
## # A tibble: 1 x 2  
##       n mean_price  
##   <int>      <dbl>  
## 1  53940      3933.
```

Your turn 7

Use summarise() to compute three statistics about the gapminder data set:

- 1. The first (min()) year in the data**
- 2. The last (max()) year in the data**
- 3. The total number of observations (n()) and the total number of unique countries in the data (n_distinct())**

Your turn 7

```
gapminder %>%  
  summarize(  
    first = min(year),  
    last = max(year),  
    n = n(),  
    n_countries = n_distinct(country)  
  )
```

```
## # A tibble: 1 x 4  
##   first last      n n_countries  
##   <dbl> <dbl> <int>      <int>  
## 1  1952  2007  1704        142
```

group_by()

```
group_by(<DATA>, <VARIABLE>)
```


group_by()

```
diamonds %>%  
  group_by(cut)
```

group_by()

```
diamonds %>%  
  group_by(cut)
```

```
## # A tibble: 53,940 x 10
```

```
## # Groups:   cut [5]
```

	carat	cut	color	clarity	depth	table	price	x	y
	<dbl>	<ord>	<ord>	<ord>	<dbl>	<dbl>	<int>	<dbl>	<dbl>
## 1	0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98
## 2	0.21	Premium	E	SI1	59.8	61	326	3.89	3.84
## 3	0.23	Good	E	VS1	56.9	65	327	4.05	4.07
## 4	0.290	Premium	I	VS2	62.4	58	334	4.2	4.23
## 5	0.31	Good	J	SI2	63.3	58	335	4.34	4.35
## 6	0.24	Very G...	J	VVS2	62.8	57	336	3.94	3.96
## 7	0.24	Very G...	I	VVS1	62.3	57	336	3.95	3.98
## 8	0.26	Very G...	H	SI1	61.9	55	337	4.07	4.11
## 9	0.22	Fair	E	VS2	65.1	61	337	3.87	3.78
## 10	0.23	Very G...	H	VS1	59.4	61	338	4	4.05

```
## # ... with 53,930 more rows, and 1 more variable: z <dbl>
```

group_by()

```
diamonds %>%  
  group_by(cut) %>%  
  summarize(n = n(), mean_price = mean(price))
```

group_by()

```
diamonds %>%  
  group_by(cut) %>%  
  summarize(n = n(), mean_price = mean(price))
```

```
## # A tibble: 5 x 3  
##   cut          n mean_price  
##   <ord>      <int>      <dbl>  
## 1 Fair       1610       4359.  
## 2 Good       4906       3929.  
## 3 Very Good 12082       3982.  
## 4 Premium   13791       4584.  
## 5 Ideal     21551       3458.
```

group_by()

```
diamonds %>%  
  group_by(cut) %>%  
  mutate(n = n(), mean_price = mean(price))
```

group_by()

```
diamonds %>%  
  group_by(cut) %>%  
  mutate(n = n(), mean_price = mean(price))
```

```
## # A tibble: 53,940 x 12
```

```
## # Groups:   cut [5]
```

	carat	cut	color	clarity	depth	table	price	x	y
	<dbl>	<ord>	<ord>	<ord>	<dbl>	<dbl>	<int>	<dbl>	<dbl>
## 1	0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98
## 2	0.21	Premium	E	SI1	59.8	61	326	3.89	3.84
## 3	0.23	Good	E	VS1	56.9	65	327	4.05	4.07
## 4	0.290	Premium	I	VS2	62.4	58	334	4.2	4.23
## 5	0.31	Good	J	SI2	63.3	58	335	4.34	4.35
## 6	0.24	Very G...	J	VVS2	62.8	57	336	3.94	3.96
## 7	0.24	Very G...	I	VVS1	62.3	57	336	3.95	3.98
## 8	0.26	Very G...	H	SI1	61.9	55	337	4.07	4.11
## 9	0.22	Fair	E	VS2	65.1	61	337	3.87	3.78
## 10	0.23	Very G...	H	VS1	59.4	61	338	4	4.05

```
## # ... with 53,930 more rows, and 3 more variables: z <dbl>,  
## #   n <int>, mean_price <dbl>
```

Your turn 8

Extract the rows where continent == "Europe". Then use group_by() to group by country. Finally, use summarize() to compute:

- 1. The total number of observations for each country in Europe**
- 2. The lowest observed life expectancy for each country**

Your turn 8

```
gapminder %>%  
  filter(continent == "Europe") %>%  
  group_by(country) %>%  
  summarize(n = n(), min_le = min(lifeExp))
```

```
## # A tibble: 30 x 3
```

```
##   country          n min_le  
##   <fct>          <int> <dbl>  
## 1 Albania         12   55.2  
## 2 Austria         12   66.8  
## 3 Belgium         12    68  
## 4 Bosnia and Herzegovina 12   53.8  
## 5 Bulgaria         12   59.6  
## 6 Croatia          12   61.2  
## 7 Czech Republic   12   66.9  
## 8 Denmark          12   70.8  
## 9 Finland          12   66.6  
## 10 France          12   67.4  
## # ... with 20 more rows
```


Your turn 9

Use grouping to calculate the mean life expectancy for each continent and year. Call the mean life expectancy variable `mean_le`. Plot the life expectancy over time (no need to change the plot code).

```
gapminder %>%  
  ----- %>%  
  ----- %>%  
  ggplot(aes(x = year, y = mean_le, col = continent)) +  
    geom_line() +  
    scale_color_manual(values = continent_colors)
```

Your turn 9

Use grouping to calculate the mean life expectancy for each continent and year. Call the mean life expectancy variable `mean_le`. Plot the life expectancy over time (no need to change the plot code).

```
gapminder %>%  
  group_by(continent, year) %>%  
  summarize(mean_le = mean(lifeExp)) %>%  
  ggplot(aes(x = year, y = mean_le, col = continent)) +  
    geom_line() +  
    scale_color_manual(values = continent_colors)
```

`mutate_if/all/at()`

`summarize_if/all/at()`

... etc!

Joining data

Use `left_join()`, `right_join()`, `full_join()`, **or** `inner_join()` **to join datasets**

Use `semi_join()` **or** `anti_join()` **to filter datasets against each other**

Resources

R for Data Science: A comprehensive but friendly introduction to the tidyverse.
Free online.

RStudio Primers: Free interactive courses in the Tidyverse

10 dplyr tips: a Twitter thread on other useful aspects of dplyr