Modeling in R and Tidying Results

linear models and broom

2019-08-15

⚠️ This is not a course in a regression

Modeling in R

Modeling in R

Modeling in R

Modeling in R

Modeling in R

lm() = **Linear Regression (OLS)**

Modeling in R

lm() = Linear Regression (OLS)

glm() = **Generalized Linear Model
(default family = Gaussian)**

Modeling in R

```
lm(price ~ carat, data = diamonds)
```

Modeling in R

```
lm(price ~ carat, data = diamonds)
```

```
##
## Call:
## lm(formula = price ~ carat, data = diamonds)
##
## Coefficients:
## (Intercept)          carat
##       -2256           7756
```

Modeling in R

```r
lm(price ~ carat, data = diamonds) %>%
  summary()
```

Modeling in R

```
## 
## Call:
## lm(formula = price ~ carat, data = diamonds)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18585.3   -804.8    -18.9    537.4  12731.7
## 
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2256.36       13.06  -172.8   <2e-16 ***
## carat        7756.43       14.07   551.4   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 1549 on 53938 degrees of freedom
## Multiple R-squared:  0.8493,    Adjusted R-squared:  0.8493
## F-statistic: 3.041e+05 on 1 and 53938 DF,  p-value: < 2.2e-16
```

Modeling in R

```
## 
## Call:
## lm(formula = price ~ carat, data = diamonds)
## 
## Residuals:
##      Min       1Q    Median       3Q      Max
## -18585.3   -804.8     -18.9    537.4  12731.7
## 
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2256.36      13.06   -172.8   <2e-16 ***
## carat        7756.43      14.07    551.4   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 1549 on 53938 degrees of freedom
## Multiple R-squared:  0.8493,	Adjusted R-squared:  0.8493
## F-statistic: 3.041e+05 on 1 and 53938 DF,  p-value: < 2.2e-16
```

broom: **tidy models**

tidy()

glance()

augment()

broom: **tidy models**

tidy() = **model coefficients**

glance()

augment()

broom: **tidy models**

tidy()

glance() = **model fit**

augment()

broom: **tidy models**

tidy()

glance()

augment() = **model predictions**

broom: **tidy models**

tidy()

glance()

augment()

NOT a core member of the tidyverse. Need to load with library(broom)

Modeling in R

```r
library(broom)
lm(price ~ carat, data = diamonds) %>%
  tidy()
```

Modeling in R

```r
library(broom)
lm(price ~ carat, data = diamonds) %>%
  tidy()
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic p.value
##   <chr>            <dbl>     <dbl>     <dbl>   <dbl>
## 1 (Intercept)     -2256.      13.1     -173.       0
## 2 carat            7756.      14.1      551.       0
```

Modeling in R

```r
lm(price ~ carat, data = diamonds) %>%
  glance()
```

Modeling in R

```r
lm(price ~ carat, data = diamonds) %>%
  glance()
```

```
## # A tibble: 1 x 11
##   r.squared adj.r.squared sigma statistic p.value    df
##       <dbl>         <dbl> <dbl>     <dbl>   <dbl> <int>
## 1     0.849         0.849 1549.   304051.       0     2
## # … with 5 more variables: logLik <dbl>, AIC <dbl>,
## #   BIC <dbl>, deviance <dbl>, df.residual <int>
```

Modeling in R

```
lm(price ~ carat, data = diamonds) %>%
  augment()
```

Modeling in R

```
lm(price ~ carat, data = diamonds) %>%
  augment()
```

```
## # A tibble: 53,940 x 9
##    price carat .fitted .se.fit .resid     .hat .sigma .cooksd
##    <int> <dbl>   <dbl>   <dbl>  <dbl>    <dbl>  <dbl>   <dbl>
##  1   326  0.23   -472.    10.4   798. 4.52e-5  1549. 6.00e-6
##  2   326  0.21   -628.    10.6   954. 4.71e-5  1549. 8.92e-6
##  3   327  0.23   -472.    10.4   799. 4.52e-5  1549. 6.02e-6
##  4   334  0.290    -7.00   9.77  341. 3.98e-5  1549. 9.66e-7
##  5   335  0.31    148.     9.57  187. 3.82e-5  1549. 2.78e-7
##  6   336  0.24   -395.    10.3   731. 4.42e-5  1549. 4.93e-6
##  7   336  0.24   -395.    10.3   731. 4.42e-5  1549. 4.93e-6
##  8   337  0.26   -240.    10.1   577. 4.24e-5  1549. 2.94e-6
##  9   337  0.22   -550.    10.5   887. 4.61e-5  1549. 7.56e-6
## 10   338  0.23   -472.    10.4   810. 4.52e-5  1549. 6.18e-6
## # … with 53,930 more rows, and 1 more variable:
## #   .std.resid <dbl>
```

Try it yourself

Work your way through the exercises. If anything in particular is giving you trouble, we'll work through it together.

Resources

R for Data Science: A comprehensive but friendly introduction to the tidyverse. Free online.

UCLA IDRE: Useful resources on modeling in R and other languages